



LLM Project

Beyond Sentiment: Turning Negative Reviews into Actionable Insights

  Korinne Angela White



Project Goals

This project creates a language model capable of efficiently analyzing negative feedback and turning it into actionable tasks, enabling businesses to quickly identify and address areas for service improvement.



Resources

Dataset:


[YELP's Full Review Dataset](#)

Pre-trained Models:


- [SamLowe/roberta-base-go_emotions](#)
- [Mistral-7B-Instruct-v0.1](#)
- [Microsoft/phi-2](#)

Process Overview

1. Use SamLowe's RoBERTa-base GoEmotions multi-class classification model on the reviews
2. Use Mistral 7B to generate constructive feedback (Gold Standard)
3. Benchmark base phi-2 to track progress
4. Quantize and fine-tune phi-2 to generate constructive feedback
5. Evaluation of the fine-tuned model



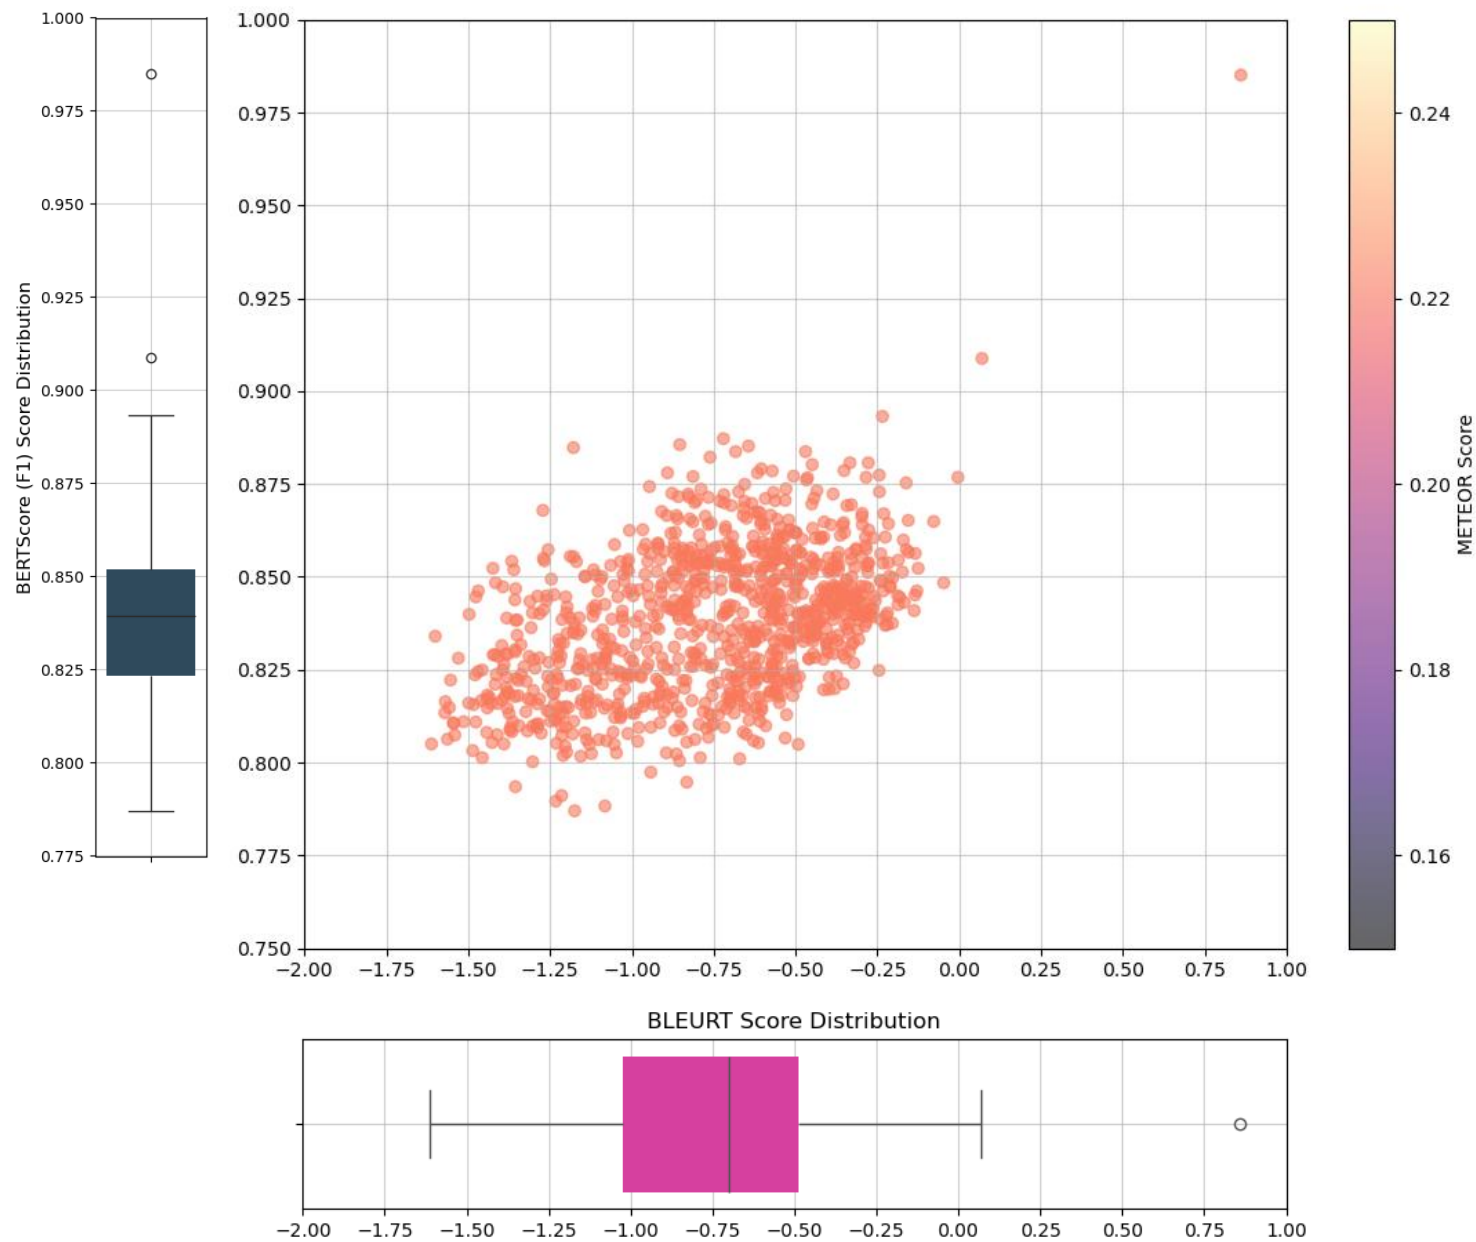
RESULTS



Metrics Overview

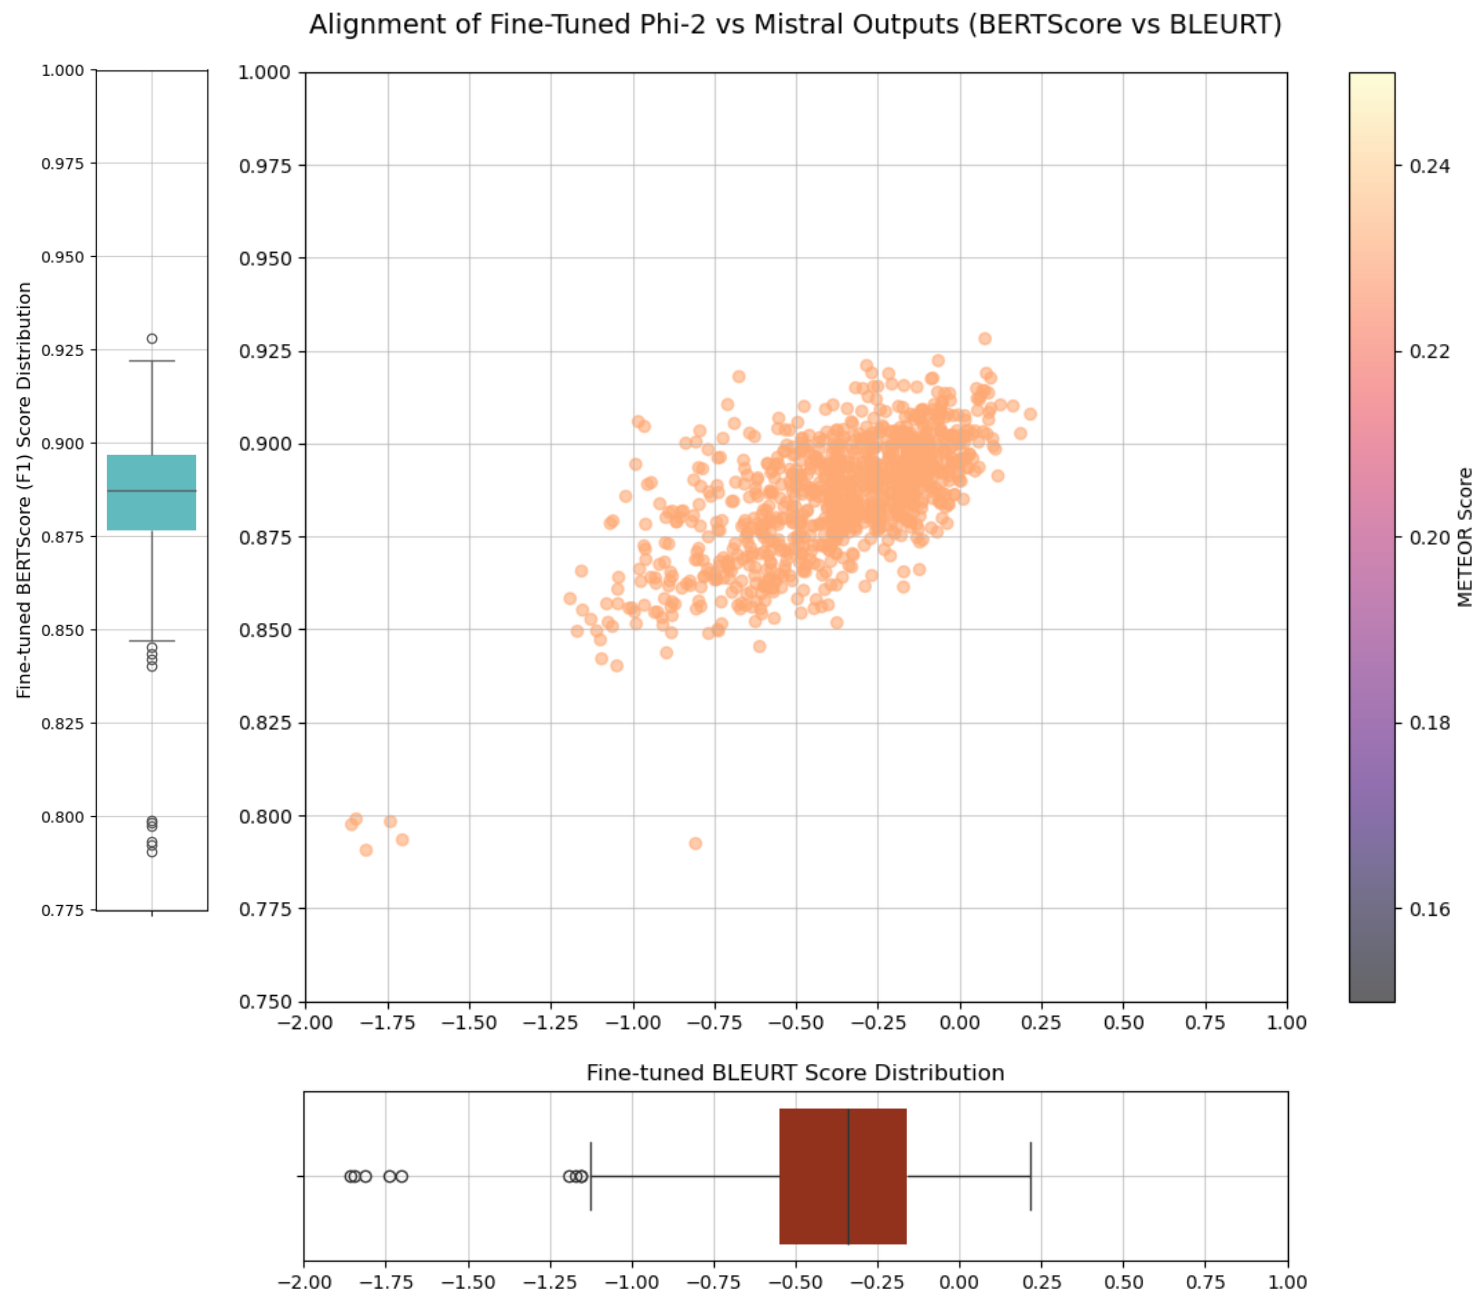
1. **BLEURT** -- Bilingual Evaluation Understudy with Representations from Transformers
2. **BERTScore** -- Bidirectional Encoder Representations from Transformers
3. **METEOR** Score -- Metric for Evaluation of Translation with Explicit Ordering

Alignment of Phi-2 vs Mistral Outputs (BERTScore vs BLEURT)



Initial
Performance

Fine-Tuned Model Performance





HYPERPARAMETERS



Hyperparameters

During Training:

- `learning_rate=2e-5`
- `metric_for_best_model="loss" + greater_is_better=False`
- `num_train_epochs=3`
- `weight_decay=0.01`
- `optim="adamw_torch_fused"`

During Generation/Inference:

- `temperature=0.5`
- `top_p=0.9`
- `no_repeat_ngram_size=3`

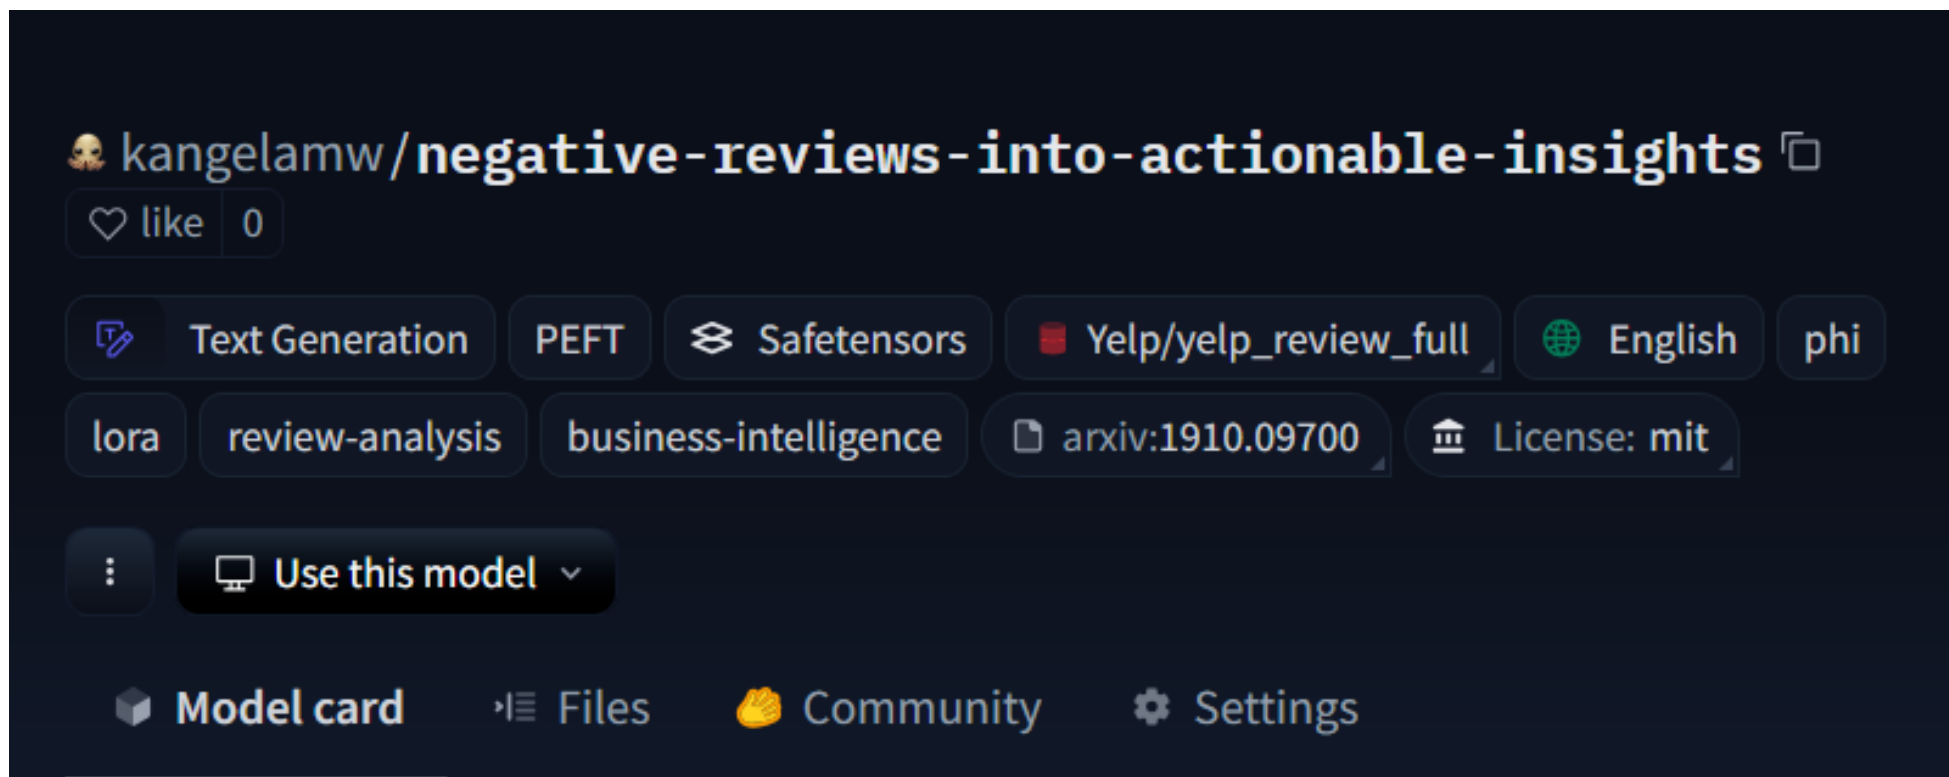




DEPLOYMENT

It's on Hugging Face!

<https://huggingface.co/kangelamw/negative-reviews-into-actionable-insights>



Technical Approach/Requirements

Quantization Method

- QLoRA 4-bit quantization

Hardware

- 12GB VRAM recommended

Library

- PEFT (Parameter-Efficient Fine-tuning)



References

Readings:

- [Enable your device for Development](#)
- [GoEmotions \(Github/Hugging Face\)](#)
- [Leveraging LLM-as-a-Judge for Automated and Scalable Evaluation](#)
- [Reliable Confidence Intervals for Information Retrieval Evaluation Using Generative A.I.](#)
- [Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena](#)
- [BLEURT: a Transfer Learning-Based Metric for Natural Language Generation](#)
- [BERTScore](#)
- [METEOR Score](#)
- [Hugging Face Inference Providers](#)
- [Hugging Face Docs \(Evaluate, PEFT/LoRA, BitsAndBytes, Trainers\)](#)

Videos:

- [LoRA & QLoRA Fine-tuning Explained In-Depth](#)
- [Simple Training with the Transformers Trainer](#)
- [The Trainer API](#)
- [Fine-Tuning Large Language Models \(LLMs\) | w/ Example Code](#)
- [What is LoRA? Low-Rank Adaptation for finetuning LLMs EXPLAINED](#)
- [What is a Context Window?](#)